

index

January 29, 2022

1 Interactions - Lab

1.1 Introduction

In this lab, you'll explore interactions in the Ames Housing dataset.

1.2 Objectives

You will be able to: - Implement interaction terms in Python using the `sklearn` and `statsmodels` packages - Interpret interaction variables in the context of a real-world problem

1.3 Build a baseline model

You'll use a couple of built-in functions, which we imported for you below:

```
[1]: from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import cross_val_score
      from sklearn.model_selection import KFold
      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
```

If you still want to build a model in the end, you can do that, but this lab will just focus on finding meaningful insights in interactions and how they can improve R^2 values.

```
[2]: regression = LinearRegression()
```

Create a baseline model which includes all the variables we selected from the Ames housing data set to predict the house prices. Then use 10-fold cross-validation and report the mean R^2 value as the baseline R^2 .

```
[11]: ames = pd.read_csv('ames.csv')

      continuous = ['LotArea', '1stFlrSF', 'GrLivArea', 'SalePrice']
      categoricals = ['BldgType', 'KitchenQual', 'SaleType', 'MSZoning', 'Street',
                     'Neighborhood']

      ames_cont = ames[continuous]
      ames_dummies = pd.get_dummies(ames[categoricals], prefix = categoricals)
```

```

data = pd.concat([ames_cont, ames_dummies], axis = 1)

Y = data["SalePrice"]
X = data.drop("SalePrice", axis = 1)

kfold = KFold(n_splits=10, shuffle=True, random_state=1)

baseline = np.mean(cross_val_score(regression, X, Y, scoring='r2', cv=kfold))

baseline

```

```

[11]:

```

	LotArea	1stFlrSF	GrLivArea	SalePrice	BldgType_1Fam	BldgType_2fmCon	\
39	6040	1152	1152	82000	0	0	
97	10921	960	960	94750	1	0	
98	10625	835	835	83000	1	0	
110	9525	1216	1855	136900	1	0	
117	8536	1125	1125	155000	1	0	
...	
1384	9060	698	1258	105000	1	0	
1423	19690	1575	2201	274970	1	0	
1448	11767	796	1346	112000	1	0	
1452	3675	1072	1072	145000	0	0	
1459	9937	1256	1256	147500	1	0	

	BldgType_Duplex	BldgType_Twnhs	BldgType_TwnhsE	KitchenQual_Ex	...	\
39	1	0	0	0	...	
97	0	0	0	0	...	
98	0	0	0	0	...	
110	0	0	0	0	...	
117	0	0	0	0	...	
...	
1384	0	0	0	0	...	
1423	0	0	0	0	...	
1448	0	0	0	0	...	
1452	0	0	1	0	...	
1459	0	0	0	0	...	

	Neighborhood_NoRidge	Neighborhood_NridgHt	Neighborhood_OldTown	\
39	0	0	0	
97	0	0	0	
98	0	0	0	
110	0	0	0	
117	0	0	0	
...	
1384	0	0	0	

1423	0	0	0
1448	0	0	0
1452	0	0	0
1459	0	0	0

	Neighborhood_SWISU	Neighborhood_Sawyer	Neighborhood_SawyerW	\
39	0	0	0	
97	0	0	0	
98	0	0	0	
110	0	0	0	
117	0	0	0	
...	
1384	0	0	0	
1423	0	0	0	
1448	0	0	0	
1452	0	0	0	
1459	0	0	0	

	Neighborhood_Somerst	Neighborhood_StoneBr	Neighborhood_Timber	\
39	0	0	0	
97	0	0	0	
98	0	0	0	
110	0	0	0	
117	0	0	0	
...	
1384	0	0	0	
1423	0	0	0	
1448	0	0	0	
1452	0	0	0	
1459	0	0	0	

	Neighborhood_Veenker
39	0
97	0
98	0
110	0
117	0
...	...
1384	0
1423	0
1448	0
1452	0
1459	0

[100 rows x 54 columns]

1.4 See how interactions improve your baseline

Next, create all possible combinations of interactions, loop over them and add them to the baseline model one by one to see how they affect the R^2 . We'll look at the 3 interactions which have the biggest effect on our R^2 , so print out the top 3 combinations.

You will create a `for` loop to loop through all the combinations of 2 predictors. You can use `combinations` from `itertools` to create a list of all the pairwise combinations. To find more info on how this is done, have a look [here](#).

Since there are so many different neighbourhoods we will exclude

```
[4]: from itertools import combinations
      combo = combinations(X.columns, 2)
```

```
[8]: # code to find top interactions by R^2 value here
      # dic = {}
      # XX = X.copy()
      # for i, (a,b) in enumerate(combo):
      #     XX[a + "*" + b] = X[a]*X[b]
      #     results = np.mean(cross_val_score(regression, XX,Y, scoring='r2',
      #                                     ↪cv=kfold))
      #     dic[np.round(results,3)] = (a,b)
      #     if i%50 == 0:
      #         print(i)
```

```
[9]: # lis = sorted(dic.keys(), reverse = True)
      # lis
```

It looks like the top interactions involve the `Neighborhood_Edwards` feature so lets add the interaction between `LotArea` and `Edwards` to our model.

We can interpret this feature as the relationship between `LotArea` and `SalePrice` when the house is in `Edwards` or not.

1.5 Visualize the Interaction

Separate all houses that are located in `Edwards` and those that are not. Run a linear regression on each population against `SalePrice`. Visualize the regression line and data points with price on the y axis and `LotArea` on the x axis.

```
[ ]:
[61]: # Visualization code here
      import seaborn as sns
      sns.set(rc={"figure.figsize":(10, 10)})

      to_pick = ["LotArea", "SalePrice"]
```

```

Edwards = data.loc[data["Neighborhood_Edwards"] == 1][to_pick]
Not_Edwards = data.loc[data["Neighborhood_Edwards"] == 0][to_pick]

model_edward = LinearRegression()
model_not_edward = LinearRegression()

X1 = np.reshape(np.log(np.array(Edwards["LotArea"])), (-1,1))
X2 = np.reshape(np.log(np.array(Not_Edwards["LotArea"])), (-1,1))

Y1 = np.log(Edwards["SalePrice"])
Y2 = Y = np.log(Not_Edwards["SalePrice"])

model_edward.fit(X1, Y1)
model_not_edward.fit(X2, Y2)

prediction_edward = model_edward.predict(X1)
prediction_not_edward = model_not_edward.predict(X2)

sns.scatterplot(x = np.log(np.array(Edwards["LotArea"])),
                y = np.array(Y1), color = "red", label = "Edwards")

sns.scatterplot(x = np.log(np.array(Not_Edwards["LotArea"])),
                y = np.array(Y2), color = "blue", label = "Not Edwards",
                alpha = 0.1);

sns.lineplot(x = np.log(np.array(Edwards["LotArea"])),
             y = prediction_edward, color = "red",
             label = "Model Prediction in Edward")

sns.lineplot(x = np.log(np.array(Not_Edwards["LotArea"])),
             y = prediction_not_edward, color = "blue",
             label = "Model Prediction Not in Edward");

```



1.6 Build a final model with interactions

Use 10-fold cross-validation to build a model using the above interaction.

```
[63]: # code here

Y = data["SalePrice"]
X = data.drop("SalePrice", axis = 1)

X["interaction"] = X["Neighborhood_Edwards"] * X["LotArea"]

kfold = KFold(n_splits=10, shuffle=True, random_state=1)
```

```
baseline_final = np.mean(cross_val_score(regression, X, Y, scoring='r2',
↪cv=kfold))

baseline_final
```

[63]: 0.8093314939294032

Our R^2 has increased considerably! Let's have a look in `statsmodels` to see if this interactions are significant.

```
[65]: # code here
import statsmodels.api as sm

X_added = sm.add_constant(X)
model = sm.OLS(Y,X_added)
results = model.fit()

print(results.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  SalePrice    R-squared:                0.835
Model:                            OLS      Adj. R-squared:            0.829
Method:                    Least Squares    F-statistic:                148.6
Date:                Sat, 29 Jan 2022      Prob (F-statistic):          0.00
Time:                02:21:35              Log-Likelihood:             -17229.
No. Observations:                1460      AIC:                        3.456e+04
Df Residuals:                    1411      BIC:                        3.482e+04
Df Model:                        48
Covariance Type:                nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const                2.082e+04    4252.352        4.896      0.000      1.25e+04
2.92e+04
LotArea                0.6108         0.103        5.916      0.000         0.408
0.813
1stFlrSF                35.0664         3.288       10.664      0.000        28.616
41.517
GrLivArea                58.1426         2.405       24.171      0.000        53.424
62.861
BldgType_1Fam          2.602e+04    2638.692         9.861      0.000      2.08e+04
3.12e+04
BldgType_2fmCon        9217.1096     5696.381         1.618      0.106     -1957.176

```

2.04e+04					
BldgType_Duplex 2165.737	-6841.6893	4591.768	-1.490	0.136	-1.58e+04
BldgType_Twnhs 668.832	-1.026e+04	5569.935	-1.842	0.066	-2.12e+04
BldgType_TwnhsE 9980.788	2679.5402	3722.000	0.720	0.472	-4621.708
KitchenQual_Ex 6.38e+04	5.641e+04	3780.490	14.921	0.000	4.9e+04
KitchenQual_Fa -1.54e+04	-2.433e+04	4554.502	-5.342	0.000	-3.33e+04
KitchenQual_Gd 6809.881	2308.8324	2294.526	1.006	0.314	-2192.217
KitchenQual_TA -9345.160	-1.357e+04	2152.934	-6.302	0.000	-1.78e+04
SaleType_COD -5228.796	-1.794e+04	6477.445	-2.769	0.006	-3.06e+04
SaleType_CWD 3.35e+04	3318.0182	1.54e+04	0.216	0.829	-2.69e+04
SaleType_Con 8.75e+04	4.498e+04	2.17e+04	2.075	0.038	2451.689
SaleType_ConLD 2.02e+04	-1510.0213	1.11e+04	-0.136	0.892	-2.33e+04
SaleType_ConLI 2.64e+04	-889.2966	1.39e+04	-0.064	0.949	-2.82e+04
SaleType_ConLw 2.03e+04	-7014.0283	1.39e+04	-0.503	0.615	-3.44e+04
SaleType_New 2.51e+04	1.434e+04	5483.489	2.614	0.009	3579.608
SaleType_Oth 2.52e+04	-9223.3553	1.76e+04	-0.525	0.600	-4.37e+04
SaleType_WD 3801.147	-5248.4845	4613.283	-1.138	0.255	-1.43e+04
MSZoning_C (all) 912.451	-1.969e+04	1.05e+04	-1.875	0.061	-4.03e+04
MSZoning_FV 3.34e+04	1.83e+04	7715.456	2.371	0.018	3161.253
MSZoning_RH 1.38e+04	-1635.0973	7891.083	-0.207	0.836	-1.71e+04
MSZoning_RL 1.73e+04	9561.0427	3920.921	2.438	0.015	1869.582
MSZoning_RM 2.3e+04	1.428e+04	4419.309	3.232	0.001	5615.470
Street_Grvl 2.94e+04	1.203e+04	8875.663	1.356	0.175	-5377.530
Street_Pave 2.08e+04	8785.4294	6111.477	1.438	0.151	-3203.130
Neighborhood_Blmngtn	1.02e+04	8807.813	1.158	0.247	-7077.770

2.75e+04					
Neighborhood_Blueste	1.821e+04	2.29e+04	0.794	0.427	-2.68e+04
6.32e+04					
Neighborhood_BrDale	-1003.6293	9697.145	-0.103	0.918	-2e+04
1.8e+04					
Neighborhood_BrkSide	-3.116e+04	5053.420	-6.166	0.000	-4.11e+04
-2.12e+04					
Neighborhood_ClearCr	-1.337e+04	6616.187	-2.020	0.044	-2.63e+04
-386.442					
Neighborhood_CollgCr	2241.7603	3289.701	0.681	0.496	-4211.471
8694.992					
Neighborhood_Crawfor	2102.0934	4845.481	0.434	0.664	-7403.028
1.16e+04					
Neighborhood_Edwards	3.228e+04	6269.634	5.149	0.000	2e+04
4.46e+04					
Neighborhood_Gilbert	-1491.7134	4183.056	-0.357	0.721	-9697.392
6713.965					
Neighborhood_IDOTRR	-4.174e+04	7493.343	-5.571	0.000	-5.64e+04
-2.7e+04					
Neighborhood_MeadowV	-1.728e+04	8868.480	-1.949	0.052	-3.47e+04
113.443					
Neighborhood_Mitchel	-9966.6701	4957.856	-2.010	0.045	-1.97e+04
-241.108					
Neighborhood_NAMes	-2.442e+04	3017.409	-8.092	0.000	-3.03e+04
-1.85e+04					
Neighborhood_NPkVill	1.652e+04	1.13e+04	1.461	0.144	-5655.498
3.87e+04					
Neighborhood_NWames	-1.636e+04	4279.043	-3.824	0.000	-2.48e+04
-7969.153					
Neighborhood_NoRidge	6.052e+04	5779.198	10.472	0.000	4.92e+04
7.19e+04					
Neighborhood_NridgHt	5.279e+04	4722.046	11.179	0.000	4.35e+04
6.21e+04					
Neighborhood_OldTown	-4.98e+04	4947.695	-10.064	0.000	-5.95e+04
-4.01e+04					
Neighborhood_SWISU	-4.812e+04	6950.579	-6.924	0.000	-6.18e+04
-3.45e+04					
Neighborhood_Sawyer	-2.517e+04	4291.324	-5.864	0.000	-3.36e+04
-1.67e+04					
Neighborhood_SawyerW	-5213.2848	4655.479	-1.120	0.263	-1.43e+04
3919.121					
Neighborhood_Somerst	1.426e+04	7448.993	1.914	0.056	-355.522
2.89e+04					
Neighborhood_StoneBr	6.415e+04	7064.721	9.081	0.000	5.03e+04
7.8e+04					
Neighborhood_Timber	6900.0704	5705.660	1.209	0.227	-4292.418
1.81e+04					
Neighborhood_Veenker	2.572e+04	9988.488	2.575	0.010	6130.720

```

4.53e+04
interaction          -7.1552      0.513   -13.959      0.000      -8.161
-6.150
=====
Omnibus:              381.039   Durbin-Watson:              1.945
Prob(Omnibus):        0.000   Jarque-Bera (JB):          3465.080
Skew:                 0.947   Prob(JB):                  0.00
Kurtosis:             10.306   Cond. No.                  5.53e+17
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.02e-24. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

What is your conclusion here?

```
[ ]: # formulate your conclusion
```

1.7 Summary

You should now understand how to include interaction effects in your model! As you can see, interactions can have a strong impact on linear regression models, and they should always be considered when you are constructing your models.